

DapR1: Notes on the Live R Session, Week 8

This week, we will be picking up where we left off the last time. We'll continue working with probability using operators. We'll also talk about writing a clean report using Markdown, including learning about in-line R coding and notation.

First, let's load the tidyverse and create our sample space. We'll also recreate our sample data from last week. Because we used the `set.seed()` function last week, we can use the same seed value to make R generate the same data we used last week.

```
library(tidyverse)
skittles <- c('red', 'green', 'yellow', 'purple', 'orange')

set.seed(2210)
s5 <- tibble(0bs = sample(skittles, 5, replace = T))
s100 <- tibble(0bs = sample(skittles, 100, replace = T))
```

Although we're working with character data in this example, you can also create sample data with numeric values. Consider an example where you want to calculate the probability of rolling an even number on a die. You could use `seq()` function, which allows you to generate sequences of numbers, to create your sample space and your events of interest:

```
dSp <- seq(1, 6)
dEv <- seq(2, 6, by = 2)
```

Joint Probability

With joint probability, we're interested in looking at the probability of multiple outcomes. For example, if we were to select a skittle from the bag, what's the likelihood we would select a yellow **OR** a purple?

Because these events are mutually exclusive, we would expect their joint probability to be the sum of their individual probabilities.

$$p(\text{yellow} \cup \text{purple}) = .20 + .20 = .40$$

Here, we'll review two ways you can specify multiple events of interest and use these methods to compute joint probability in your sample.

The AND and OR operators

Two operators that may be used for specifying multiple conditions are `&` and `|`. As you might expect, `&` refers to 'and'. It allows you to specify values that meet *all* given conditions. The `|` operator stands for 'or' and allows you to specify values that meet *any* of the given conditions.

```
successes <- s100 == 'yellow' | s100 == 'purple'
failures <- s100 != 'yellow' & s100 != 'purple'

sum(successes)/(sum(successes) + sum(failures))
```

```
## [1] 0.4
```

`%in%` or `is.element()`

You can get the same results using the `%in%` operator. This operator can be used to determine whether a given element (or values from a vector of elements) is found in a dataset or vector. Also, note the use of another operator. The `!` operator stands for ‘not’ and allows you to select values that do not meet a specified condition.

```
events <- c('purple', 'yellow')
successes <- s100$Obs %in% events
failures <- !s100$Obs %in% events

sum(successes)/(sum(successes) + sum(failures))
```

```
## [1] 0.4
```

You can also use the `is.element()` function in the same way. *is.element(x, y)* is identical to *x %in% y*

```
is.element(s100$Obs, events)
```

```
## [1] FALSE TRUE TRUE FALSE FALSE TRUE FALSE TRUE TRUE FALSE TRUE TRUE
## [13] FALSE TRUE FALSE FALSE FALSE TRUE FALSE TRUE TRUE FALSE TRUE FALSE
## [25] FALSE FALSE FALSE TRUE FALSE TRUE FALSE TRUE TRUE TRUE TRUE FALSE
## [37] TRUE FALSE FALSE TRUE FALSE FALSE FALSE FALSE FALSE TRUE TRUE TRUE
## [49] FALSE FALSE FALSE TRUE TRUE FALSE TRUE FALSE TRUE FALSE TRUE FALSE
## [61] FALSE TRUE FALSE FALSE FALSE TRUE TRUE TRUE FALSE FALSE TRUE TRUE
## [73] TRUE TRUE FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE
## [85] FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE TRUE FALSE FALSE FALSE
## [97] FALSE TRUE FALSE FALSE
```

```
successes <- is.element(s100$Obs, events)
failures <- !is.element(s100$Obs, events)
sum(successes)/(sum(successes) + sum(failures))
```

```
## [1] 0.4
```

Now, let’s compute the probability of selecting a yellow **AND** a purple skittle in two selections.

To do this, we would use we use the following formula:

$$p(\text{yellow} \cap \text{purple}) = .20 * .20 = .04$$

Let’s check to see if our data follow the expected result. To do this, we’ll need to add a new column that represents the second skittle selection:

```
s100$Obs2 <- sample(skittles, 100, replace = T)
head(s100)
```

```
## # A tibble: 6 x 2
##   Obs   Obs2
##   <chr> <chr>
## 1 green orange
## 2 yellow orange
## 3 yellow orange
## 4 orange purple
## 5 red     purple
## 6 yellow purple
```

```

successes <- (s100$Obs=='purple'&s100$Obs2=='yellow')|(s100$Obs=='yellow'&s100$Obs2=='purple')
failures <- !successes

sum(successes)/(sum(successes) + sum(failures))

```

```
## [1] 0.05
```

You can also look at the proportion table for multiple events:

```

s100 %>%
  table %>%
  prop.table()

```

```

##           Obs2
## Obs      green orange purple  red yellow
## green    0.01   0.05   0.01 0.05   0.06
## orange   0.04   0.05   0.06 0.02   0.03
## purple   0.03   0.08   0.06 0.02   0.02
## red      0.06   0.04   0.05 0.04   0.03
## yellow   0.06   0.06   0.03 0.03   0.01

```

Note that you can sum the appropriate locations in the proportion table to get the corresponding probability.

```

pTab <- s100 %>%
  table() %>%
  prop.table()

pTab['purple', 'yellow'] + pTab['yellow', 'purple']

```

```
## [1] 0.05
```

You can also compute the proportions for each row or each column using the *margin* argument:

```

# margin = 1: proportion within each row
pTab <- s100 %>%
  table %>%
  prop.table(margin = 1)

pTab

```

```

##           Obs2
## Obs      green      orange      purple      red      yellow
## green  0.05555556 0.27777778 0.05555556 0.27777778 0.33333333
## orange 0.20000000 0.25000000 0.30000000 0.10000000 0.15000000
## purple 0.14285714 0.38095238 0.28571429 0.09523810 0.09523810
## red    0.27272727 0.18181818 0.22727273 0.18181818 0.13636364
## yellow 0.31578947 0.31578947 0.15789474 0.15789474 0.05263158

```

```
sum(pTab[1,])
```

```
## [1] 1
```

```
# margin = 2: proportion within each column
```

```
pTab <- s100 %>%
```

```
  table %>%
```

```
  prop.table(margin = 2)
```

```
pTab
```

```
##      Obs2
```

```
## Obs      green      orange      purple      red      yellow
```

```
## green 0.05000000 0.17857143 0.04761905 0.31250000 0.40000000
```

```
## orange 0.20000000 0.17857143 0.28571429 0.12500000 0.20000000
```

```
## purple 0.15000000 0.28571429 0.28571429 0.12500000 0.13333333
```

```
## red    0.30000000 0.14285714 0.23809524 0.25000000 0.20000000
```

```
## yellow 0.30000000 0.21428571 0.14285714 0.18750000 0.06666667
```

```
sum(pTab[,1])
```

```
## [1] 1
```

However, the table outputs in R are not appropriate for a formal report. If you want to include a table in your report, you should use the `kbl()` function to produce a formal table to display your data:

```
library(kableExtra)
```

```
pTab %>%
```

```
  kbl(booktabs = T, digits = 2, caption = 'Skittles Experiment Data') %>%
```

```
  kable_styling(latex_options = c('hold_position', 'striped')) %>%
```

```
  column_spec(1, bold = T) %>%
```

```
  row_spec(0, bold = T)
```

Table 1: Skittles Experiment Data

	green	orange	purple	red	yellow
green	0.05	0.18	0.05	0.31	0.40
orange	0.20	0.18	0.29	0.12	0.20
purple	0.15	0.29	0.29	0.12	0.13
red	0.30	0.14	0.24	0.25	0.20
yellow	0.30	0.21	0.14	0.19	0.07

Conditional Probability Data & Write-Up

Now let's import some data that we can use to demonstrate conditional probability. We'll also write up our results neatly and talk about in-line R coding.

Imagine that we want to investigate the relationship between med school acceptance and academic performance. Specifically we will look at the likelihood of acceptance (*yes* or *no*) is related to having higher marks (e.g., above the 75th percentile or greater in our sample; *high* or...less high. We'll just say *low* for simplicity's sake).

```
dat <- read.csv('https://uoepsy.github.io/data/MedGPA.csv')
summary(dat)
```

```
##      Accept      Acceptance      Sex      BCPM
## Length:55      Min.    :0.0000 Length:55      Min.    :2.410
## Class :character 1st Qu.:0.0000 Class :character 1st Qu.:3.260
## Mode  :character Median :1.0000 Mode  :character Median :3.530
##                      Mean  :0.5455 Mean  :3.501
##                      3rd Qu.:1.0000 3rd Qu.:3.755
##                      Max.   :1.0000 Max.   :4.000
##
##      GPA      VR      PS      WS
## Min.    :2.720 Min.    : 6.000 Min.    : 5.000 Min.    : 4.000
## 1st Qu.:3.375 1st Qu.: 8.000 1st Qu.: 9.000 1st Qu.: 6.000
## Median :3.580 Median :10.000 Median :10.000 Median : 8.000
## Mean   :3.553 Mean   : 9.764 Mean   : 9.709 Mean   : 7.148
## 3rd Qu.:3.770 3rd Qu.:11.000 3rd Qu.:10.500 3rd Qu.: 8.000
## Max.   :3.970 Max.   :13.000 Max.   :14.000 Max.   :10.000
##                      NA's    :1
##
##      BS      MCAT      Apps
## Min.    : 6.000 Min.    :18.00 Min.    : 1.000
## 1st Qu.: 9.000 1st Qu.:34.00 1st Qu.: 5.000
## Median :10.000 Median :36.00 Median : 7.000
## Mean   : 9.782 Mean   :36.27 Mean   : 8.364
## 3rd Qu.:11.000 3rd Qu.:39.00 3rd Qu.:11.000
## Max.   :14.000 Max.   :48.00 Max.   :24.000
##
```

```
dat$acceptChar <- ifelse(dat$Acceptance == '0', 'Rejected', 'Accepted')
dat$GPASplit <- ifelse(dat$GPA >= quantile(dat$GPA)['75%'], 'High', 'Low')
```

```
pTab <- table(dat$GPASplit, dat$acceptChar) %>%
  prop.table()
```

```
pTab
```

```
##
##      Accepted Rejected
## High 0.25454545 0.01818182
## Low  0.29090909 0.43636364
```

Remember from yesterday's lecture that we can calculate conditional probability using the following formula:

$$p(A|B) = \frac{p(A \cap B)}{p(B)}$$

Let's compare the probability of being accepted given that marks are high with the probability given that marks are low. In other words, we are comparing an outcome of event A across different levels of event B.

If these events were unrelated, we would expect the probability of event A to be generally consistent across both levels of B. To do this, we can plug our variables into the formula above:

$$p(\text{Accepted}|\text{High}) = \frac{p(\text{Accepted} \cap \text{High})}{p(\text{High})}$$

$$p(\text{Accepted}|\text{Low}) = \frac{p(\text{Accepted} \cap \text{Low})}{p(\text{Low})}$$

Remember, we can pull specific values from the probability table we've created:

```
pTab['High', 'Accepted']/(sum(pTab['High',]))
```

```
## [1] 0.9333333
```

```
pTab['Low', 'Accepted']/(sum(pTab['Low',]))
```

```
## [1] 0.4
```

Here, we see that the values are quite different, which indicates that the probability of acceptance is likely to be related to overall marks (shocking!). Now let's show an example of how to write up these results.

EXAMPLE WRITE-UP

In this experiment, we collected data on medical school acceptance rates and school performance from 55 participants. Specifically, we investigated whether acceptance to medical school was related to overall performance as measured by GPA (High/Low). We calculated the proportion of participants who fell into each category (see Table 2).

Table 2: Medical School Acceptance by GPA

	Accepted	Rejected
High	0.25	0.02
Low	0.29	0.44

We computed the probability of acceptance at both levels of GPA. The probability of acceptance given a high GPA was 0.93. The probability of acceptance given a low GPA was only 0.4. This indicates that the likelihood of being accepted changes at different levels of GPA.

The differences in the probability of acceptance across levels of GPA indicates a possible relationship between school performance and acceptance to medical school.